

How to apply to college

Max Kapur

February 28, 2022

Abstract

This paper considers the maximization of the expected maximum value of a portfolio of random variables subject to a budget constraint. We refer to this as the optimal college application problem. When each variable's cost, or each college's application fee, is identical, we show that the optimal portfolios are nested in the budget constraint, yielding an exact polynomial-time algorithm. When colleges differ in their application fees, we show that the problem is NP-complete. We provide two dynamic programs for this more general setup. The first produces an exact solution in pseudopolynomial time, and the second yields a fully polynomial-time approximation scheme.

Contents

1	Introduction	3
1.1	Structure of this paper	4
2	Notation and preliminaries	4
3	Homogeneous application costs	5
3.1	Approximation properties of a naïve solution	6
3.2	The nestedness property	7
3.3	Polynomial-time solution	8
3.4	Properties of the optimal portfolios	10
4	Heterogeneous application costs	10
4.1	NP-completeness	11
4.2	Pseudopolynomial-time dynamic program	12
4.3	Fully polynomial-time approximation scheme	13
5	(WIP) Numerical experiments	16
6	References	17

1 Introduction

This paper considers the optimization problem

$$\begin{aligned} & \text{maximize} && \mathbb{E} [\max\{t_0, \max\{t_j Z_j : j \in \mathcal{X}\}\}] \\ & \text{subject to} && \mathcal{X} \subseteq \mathcal{C}, \quad \sum_{j \in \mathcal{X}} g_j \leq H \end{aligned} \tag{1}$$

where $\mathcal{C} = \{1 \dots m\}$ is an index set, H is a budget parameter, and for $j = 1 \dots m$, $g_j > 0$ is a cost parameter, Z_j is a random, independent Bernoulli variable with probability f_j , and $t_j > 0$ is a utility parameter. Without loss of generality, we assume that the t_j -values are weakly increasing in j , that $0 \leq t_0 \leq t_1$, and that each $g_j \leq H$.

We refer to this problem as the *optimal college application portfolio*, as follows: Consider a college admissions market with m college. The j th college is named c_j . Consider a single student, and let f_j denote her probability of being admitted to c_j if she applies, so that Z_j equals one if she is admitted and zero if not. It is appropriate to assume that the Z_j are probabilistically independent as long as f_j are probabilities estimated specifically for this student (as opposed to generic acceptance rates). Each t_j -value indicates the utility associated with attending c_j , and t_0 gives the utility associated with not attending college, an option named c_0 . The application fee for c_j is g_j , and H represents the student's total budget to spend on application fees. Then the student's objective is to maximize the expected utility associated with the best school she gets into, and her optimal college application strategy is given by the solution \mathcal{X} to the problem above, where \mathcal{X} represents the set of schools to which the student applies.

As Chao (2014) remarked, this represents a somewhat subtle portfolio optimization problem. Traditional portfolio optimization models weigh the expected profit across all assets against a risk term, yielding a concave maximization problem with linear constraints (Meucci 2005). But college applicants maximize the observed value of their *best* asset: If a student is admitted to her j th choice, then she is indifferent as to whether she gets into her $(j + 1)$ th choice. As a result, the valuation function that students maximize is *convex* in the expected utility associated with individual applications. Risk management is implicit in the college application problem because, in a typical admissions market, college preferability is negatively correlated with competitiveness. That is, students negotiate a tradeoff between highly attractive, selective “reach schools” and less preferable “safety schools” where admission is a safer bet. Finally, the combinatorial nature of the college application problem makes it difficult to solve using the gradient-based techniques used in continuous portfolio optimization. Chao estimated her model (which considers application as a *cost* rather than a constraint) by clustering the schools so that $m = 8$, a scale at which enumeration is possible. Our study pursues a more general solution.

Our study takes special interest in the validity of greedy solution algorithms, such as those that iteratively add the asset that yields the greatest increase in the objective function until the budget is exhausted. For certain classes of optimization problems, such as maximizing a submodular set function over a cardinality constraint, the greedy algorithm is known to be a good approximate solution and exact under certain additional assumptions (Fisher et al. 1978). For other problems, most notably the binary knapsack problem, the most obvious greedy algorithm can be made to perform arbitrarily poorly (Vazirani 2001).

We show analogous results for the college application problem: In the special case where each $g_j = 1$, the optimal portfolio satisfy a nestedness property that is equivalent to the validity of the greedy algorithm. As Rozanov and Tamir (2020) remark, the nestedness property is useful not only for obtaining solution algorithms, but in the implementation of an optimal policy under uncertain information. In the United States, many college applications are due at the beginning of November, and it is typical for students to begin working on their applications during the prior summer because colleges increasingly expect students to tailor their essays to the target school. Therefore, students may not know how many schools they can afford to apply to (in terms of both time and monetary costs) until late October. The nestedness property implies that students can prioritize their application strategy by working on applications in the order that they appear in the sequence of optimal application portfolios.

Unfortunately, the nestedness property does not hold in the general case, nor does the greedy algorithm offers any performance guarantee. We offer an alternative approximation scheme that produces a $(1 - \varepsilon)$ -optimal solution in fully polynomial time.

1.1 Structure of this paper

Section 2 introduces some additional notation and assumptions that can be imposed with trivial loss of generality.

In section 3, we solve the special case where each $g_j = 1$ and H is an integer $h \leq m$. This case mirrors the centralized college application process in Korea, where there is no application fee, but students are allowed to apply to only three schools during the main admissions cycle. We show that an intuitive heuristic is in fact a $1/h$ -approximation algorithm. Then, we show that the optimal portfolios are nested in the budget constraint, which yields an exact algorithm that runs in $O(hm)$ -time.

In section 4, we turn to the scenario in which colleges differ in their application fees. We show that the decision form of the portfolio optimization problem is NP-complete through a reduction from the binary knapsack problem. We provide two dynamic programs for this more general setup. The first iterates on total expenditures and produces an exact solution in pseudopolynomial time, namely $O(Hm + m \log m)$. The second iterates on truncated portfolio valuations and yields a fully polynomial-time approximation scheme that produces a $(1 - \varepsilon)$ -optimal solution in $O(m^3/\varepsilon)$ time.

In section 5, we present the results of computational experiments that confirm the validity and time complexity results established in the previous two sections.

2 Notation and preliminaries

Before discussing the solution algorithms, we will introduce some additional notation and a few preliminary results that will come in handy.

We refer to the set $\mathcal{X} \subseteq \mathcal{C}$ of schools to which a student applies is called her *application portfolio*. The expected utility the student receives from \mathcal{X} is called its *valuation*.

Definition 1 (Portfolio valuation function). $v(\mathcal{X}) = \mathbb{E} [\max\{t_0, \max\{t_j Z_j : j \in \mathcal{X}\}\}]$.

It is helpful to define the random variable $X = \max\{t_j Z_j : j \in \mathcal{X}\}$ as the utility achieved by the schools in the portfolio, so that when $t_0 = 0$, $v(\mathcal{X}) = \mathbb{E}[X]$. Similar pairs of variables with italic and script names such as \mathcal{Y}_h and Y_h carry an analogous meaning.

Given an application portfolio, let $p_j(\mathcal{X})$ denote the probability that the student attends c_j . This occurs if and only if she *applies* to c_j , is *admitted* to c_j , and is *rejected* from any school she prefers to c_j ; that is, any school with higher index. Hence, for $j = 0 \dots m$,

$$p_j(\mathcal{X}) = \begin{cases} f_j \prod_{\substack{i \in \mathcal{X}: \\ i > j}} (1 - f_i), & j \in \{0\} \cup \mathcal{X} \\ 0, & \text{otherwise} \end{cases} \quad (2)$$

where the empty product equals one. The following proposition follows immediately.

Proposition 1 (Closed form of portfolio valuation function).

$$v(\mathcal{X}) = \sum_{j=0}^m t_j p_j(\mathcal{X}) = \sum_{j \in \{0\} \cup \mathcal{X}} \left(t_j f_j \prod_{\substack{i \in \mathcal{X}: \\ i > j}} (1 - f_i) \right) \quad (3)$$

Next, we show that without loss of generality, we may assume that $t_0 = 0$ (or any constant less than t_1).

Theorem 1. *Let $\bar{t}_j = t_j - \gamma$ for $j = 0 \dots m$. Then $v(\mathcal{X}; \bar{t}_j) = v(\mathcal{X}; t_j) - \gamma$ regardless of \mathcal{X} .*

Proof. By definition, $\sum_{j=0}^m p_j(\mathcal{X}) = \sum_{j \in \{0\} \cup \mathcal{X}} p_j(\mathcal{X}) = 1$. Therefore

$$v(\mathcal{X}; \bar{t}_j) = \sum_{j \in \{0\} \cup \mathcal{X}} \bar{t}_j p_j(\mathcal{X}) = \sum_{j \in \{0\} \cup \mathcal{X}} (t_j - \gamma) p_j(\mathcal{X}) \quad (4)$$

$$= \sum_{j \in \{0\} \cup \mathcal{X}} t_j p_j(\mathcal{X}) - \gamma = v(\mathcal{X}; t_j) - \gamma \quad (5)$$

which completes the proof. □

3 Homogeneous application costs

In this section, we derive a polynomial-time algorithm for the special case in which $g_j = 1$ and H is a constant $h \leq m$. This case is similar to the centralized college admissions process in Korea, where there is no application fee, but by law, students are allowed to apply to no more than h schools. (In the Korean case, $m = 202$ and $h = 3$.) Applying Theorem 1, we assume that $t_0 = 0$ unless otherwise noted. Throughout this section, we will call the applicant Alma, and refer to the corresponding optimization problem as Alma's problem.

Definition 2 (Alma's problem). Alma's optimal college application portfolio is given by the

solution to the following combinatorial optimization problem:

$$\begin{aligned} & \text{maximize} \quad v(\mathcal{X}) = \sum_{j \in \mathcal{X}} \left(t_j f_j \prod_{\substack{i \in \mathcal{X}: \\ i > j}} (1 - f_i) \right) \\ & \text{subject to} \quad \mathcal{X} \subseteq \mathcal{C}, \quad |\mathcal{X}| \leq h \end{aligned} \tag{6}$$

3.1 Approximation properties of a naïve solution

The expected utility associated with a single school c_j is simply $E[t_j Z_j] = t_j f_j$. It is therefore tempting to adopt the following strategy, which turns out to be inoptimal.

Definition 3 (Naïve algorithm for Alma's problem). Apply to the h schools having the highest expected utility $t_j f_j$.

The basic error of this algorithm is that it maximizes $E[\sum t_j Z_j]$ instead of $E[\max\{t_j Z_j\}]$. The latter is what Alma is truly concerned with, since in the end she can attend only one school. The following example shows that the naïve algorithm can produce a suboptimal solution.

Example 1. Suppose $m = 3$, $h = 2$, $t = (70, 80, 90)$, and $f = (0.4, 0.4, 0.3)$. Then the naïve algorithm picks $\mathcal{T} = \{1, 2\}$ with $v(\mathcal{T}) = 70(0.4)(1 - 0.4) + 80(0.4) = 48.8$. But $\mathcal{X} = \{2, 3\}$ with $v(\mathcal{X}) = 80(0.4)(1 - 0.3) + 90(0.3) = 49.4$ is the optimal solution.

In fact, the naïve algorithm is a $(1/h)$ -approximation algorithm for Alma's problem, as expressed in the following theorem.

Theorem 2. When the application limit is h , let \mathcal{X}_h denote the optimal portfolio, and \mathcal{T}_h the set of the h schools having the largest values of $t_j f_j$. Then $v(\mathcal{T}_h)/v(\mathcal{X}_h) \geq 1/h$.

Proof. Because \mathcal{T}_h maximizes the quantity $E[\sum_{j \in \mathcal{T}_h} \{t_j Z_j\}]$, we have

$$\begin{aligned} v(\mathcal{X}_h) &= E[\max_{j \in \mathcal{X}_h} \{t_j Z_j\}] \leq E[\sum_{j \in \mathcal{X}_h} \{t_j Z_j\}] \leq E[\sum_{j \in \mathcal{T}_h} \{t_j Z_j\}] \\ &= h E[\frac{1}{h} \sum_{j \in \mathcal{T}_h} \{t_j Z_j\}] \leq h E[\max_{j \in \mathcal{T}_h} \{t_j Z_j\}] = h v(\mathcal{T}_h) \end{aligned} \tag{7}$$

where the final inequality follows from the concavity of the $\max\{\}$ operator. \square

The following example establishes the tightness of the approximation factor.

Example 2. Pick any h and let $m = 2h$. For a small constant $\varepsilon \in (0, 1)$, let

$$\begin{aligned} t &= \left(\underbrace{1, \dots, 1}_h, \underbrace{\varepsilon^{-1}, \varepsilon^{-2}, \dots, \varepsilon^{-(h-1)}, \varepsilon^{-h}}_h \right) \\ \text{and } f &= \left(\underbrace{1, \dots, 1}_h, \underbrace{\varepsilon^1, \varepsilon^2, \dots, \varepsilon^{h-1}, \varepsilon^h}_h \right) \end{aligned}$$

Since all $t_j f_j = 1$, the naïve algorithm can choose $\mathcal{T}_h = \{1, \dots, h\}$, with $v(\mathcal{T}_h) = 1$. But the optimal solution is $\mathcal{X}_h = \{h+1, \dots, m\}$, with

$$v(\mathcal{X}_h) = \sum_{j=h+1}^m \left(t_j f_j \prod_{j'=j+1}^m (1 - f_{j'}) \right) = \sum_{j=1}^h (1 - \epsilon)^j \approx h.$$

Thus, as ϵ approaches zero, we have $v(\mathcal{T}_h)/v(\mathcal{X}_h) \rightarrow 1/h$. (The optimality of \mathcal{X}_h follows from the fact that it achieves the upper bound of Theorem 7.)

Hope is not lost. We can still find the optimal solution in time polynomial in h and m , as we will now show.

3.2 The nestedness property

It turns out that the solution to Alma's problem possesses a special structure: An optimal portfolio of size $h+1$ includes an optimal portfolio of size h as a subset.

Theorem 3 (Nestedness of optimal application portfolios). *There exists a sequence of portfolios $\{\mathcal{X}_h\}_{h=1}^m$ satisfying the nestedness relation*

$$\mathcal{X}_1 \subset \mathcal{X}_2 \subset \dots \subset \mathcal{X}_m. \quad (8)$$

such that each \mathcal{X}_h is an optimal application portfolio when the application limit is h .

Proof. By induction on h . Applying Theorem 1, we assume that $t_0 = 0$.

(Base case.) First, we will show that $\mathcal{X}_1 \subset \mathcal{X}_2$. To get a contradiction, suppose that the optima are $\mathcal{X}_1 = \{j\}$ and $\mathcal{X}_2 = \{k, l\}$, where we may assume that $t_k \leq t_l$. Optimality requires that

$$v(\mathcal{X}_1) = f_j t_j > v(\{k\}) = f_k t_k \quad (9)$$

and

$$v(\mathcal{X}_2) = f_k(1 - f_l)t_k + f_l t_l > v(\{j, l\}) \quad (10)$$

$$= f_j(1 - f_l)t_j + (1 - f_j)f_l t_l + f_j f_l \max\{t_j, t_l\} \quad (11)$$

$$\geq f_j(1 - f_l)t_j + (1 - f_j)f_l t_l + f_j f_l t_l \quad (12)$$

$$= f_j(1 - f_l)t_j + f_l t_l \quad (13)$$

$$\geq f_k(1 - f_l)t_k + f_l t_l = v(\mathcal{X}_2) \quad (14)$$

which is a contradiction.

(Inductive step.) Assume that $\mathcal{X}_1 \subset \dots \subset \mathcal{X}_h$, and we will show $\mathcal{X}_h \subset \mathcal{X}_{h+1}$. Let $k = \arg \max\{t_k : k \in \mathcal{X}_{h+1}\}$ and write $\mathcal{X}_{h+1} = \mathcal{Y}_h \cup \{k\}$.

Suppose $k \notin \mathcal{X}_h$. To get a contradiction, assume that $v(\mathcal{Y}_h) < v(\mathcal{X}_h)$. Then

$$\begin{aligned} v(\mathcal{X}_{h+1}) &= v(\mathcal{Y}_h \cup \{k\}) \\ &= (1 - f_k)v(\mathcal{Y}_h) + f_k t_k \\ &< (1 - f_k)v(\mathcal{X}_h) + f_k E[\max\{t_k, X_h\}] \\ &= v(\mathcal{X}_h \cup \{k\}) \end{aligned} \tag{15}$$

contradicts the optimality of \mathcal{X}_{h+1} .

Now suppose that $k \in \mathcal{X}_h$. We can write $\mathcal{X}_h = \mathcal{Y}_{h-1} \cup \{k\}$, where \mathcal{Y}_{h-1} is some portfolio of size $h - 1$. It suffices to show that $\mathcal{Y}_{h-1} \subset \mathcal{Y}_h$. By definition, \mathcal{Y}_{h-1} (respectively, \mathcal{Y}_h) maximizes the function $v(\mathcal{Y} \cup \{k\})$ over portfolios of size $h - 1$ (respectively, h) that do not include k . That is, \mathcal{Y}_{h-1} and \mathcal{Y}_h are the optimal *complements* to the singleton portfolio $\{k\}$.

We will use the function $w(\mathcal{Y})$ to grade portfolios $\mathcal{Y} \subseteq \mathcal{C} \setminus \{k\}$ according to how well they complement $\{k\}$. To construct $w(\mathcal{Y})$, let \tilde{t}_j denote the expected utility Alma receives from school c_j given that she has been admitted to c_j and applied to c_k . For $j < k$, including $j = 0$, this is $\tilde{t}_j = t_j(1 - f_k) + t_k f_k$; for $j > k$, this is $\tilde{t}_j = t_j$. This means that

$$v(\mathcal{Y} \cup \{k\}) = \sum_{j \in \{0\} \cup \mathcal{Y}} \tilde{t}_j p_j(\mathcal{Y}). \tag{16}$$

The transformation to \tilde{t} does not change the order of the t_j -values. Therefore, the expression on the right side of (16) is itself a portfolio valuation function. In the corresponding market, t is replaced by \tilde{t} and \mathcal{C} is replaced by $\mathcal{C} \setminus \{k\}$. Now, we obtain $w(\mathcal{Y})$ through one more transformation: Define $\bar{t}_j = \tilde{t}_j - \tilde{t}_0$ so that $\bar{t}_0 = 0$ and let

$$w(\mathcal{Y}) = \sum_{j \in \{0\} \cup \mathcal{Y}} \bar{t}_j p_j(\mathcal{Y}) = \sum_{j \in \{0\} \cup \mathcal{Y}} \tilde{t}_j p_j(\mathcal{Y}) - \tilde{t}_0 = v(\mathcal{Y} \cup \{k\}) - t_k f_k \tag{17}$$

where the second equality follows from Theorem 1. This identity says that the optimal complements to $\{k\}$, given by \mathcal{Y}_{h-1} and \mathcal{Y}_h , are themselves optimal portfolios of size $h - 1$ and h for the market whose objective function is $w(\mathcal{Y})$. Since $\bar{t}_0 = 0$ in the latter market, the inductive hypothesis implies that $\mathcal{Y}_{h-1} \subset \mathcal{Y}_h$, which completes the proof.¹ \square

3.3 Polynomial-time solution

Applying the result above yields an efficient algorithm for the optimal portfolio: Start with the empty set and add schools one at a time, maximizing $v(\mathcal{X} \cup \{k\})$ at each addition. Sorting t is $O(m \log m)$. At each of the h iterations, there are $O(m)$ candidates for k , and computing $v(\mathcal{X} \cup \{k\})$ is $O(h)$ using (3); therefore, the time complexity of this algorithm is $O(h^2 m + m \log m)$.

We reduce the computation time to $O(hm)$ by taking advantage of the transformation from the inductive step in the proof of Theorem 3. Once school k is added to \mathcal{X} , we remove it from the set $\mathcal{C} \setminus \mathcal{X}$ of candidates, and update the t_j -values of the remaining schools according to the

¹We thank Yim Seho for discovering this critical transformation.

following transformation:

$$\bar{t}_j = \begin{cases} t_j(1 - f_k), & t_j \leq t_k \\ t_j - t_k f_k, & t_j > t_k \end{cases} \quad (18)$$

It is easy to verify that this is the composition of the two transformations (from t to \tilde{t} , and from \tilde{t} to \bar{t}) given in the proof. Now, the *next* school added must be the optimal singleton portfolio in the modified market. But the optimal singleton portfolio consists simply of the school with the highest value of $f_j \bar{t}_j$. Therefore, by updating the t_j -values at each iteration according to (18), we eliminate the need to compute $v(\mathcal{X})$ entirely. Moreover, this algorithm does not require the schools to be indexed in ascending order by t_j , which removes the $O(m \log m)$ sorting cost.

The algorithm below outputs a list X of the h schools to which Alma should apply. The schools appear in the order of entry such that when the algorithm is run with $h = m$, the optimal portfolio of size h is given by $\mathcal{X}_h = \{X[1], \dots, X[h]\}$. The entries of the list V give the valuation thereof.

Algorithm 1: Optimal portfolio algorithm for Alma's problem.

Data: Utility values $t \in [0, \infty)^m$, admissions probabilities $f \in [0, 1]^m$, application limit $h \leq m$.

```

1  $\mathcal{C} \leftarrow \{1 \dots m\}$ ;
2  $X, V \leftarrow$  empty lists;
3 for  $i = 1 \dots h$  do
4    $k \leftarrow \arg \max_{j \in \mathcal{C}} \{f_j t_j\}$ ;
5    $\mathcal{C} \leftarrow \mathcal{C} \setminus \{k\}$ ;
6    $\text{append!}(X, k)$ ;
7   if  $i = 1$  then  $\text{append!}(V, f_k t_k)$  else  $\text{append!}(V, V[i - 1] + f_k t_k)$ ;
8   for  $j \in \mathcal{C}$  do
9     if  $t_j \leq t_k$  then  $t_j \leftarrow t_j(1 - f_k)$  else  $t_j \leftarrow t_j - f_k t_k$ ;
10  end
11 end
12 return  $X, V$ 
```

Theorem 4 (Validity of Algorithm 1). *Algorithm 1 produces an optimal application portfolio for Alma's problem in $O(hm)$ time.*

Proof. Optimality follows from the proof of Theorem 3. Suppose \mathcal{C} is stored as a list. Then at each of the h iterations of the main loop, finding the top school costs $O(m)$, and the t_j -values of the remaining $O(m)$ schools are each updated in unit time. Therefore, the overall time complexity is $O(hm)$. \square

In our numerical experiments, we found it effective to store \mathcal{C} as a binary max heap rather than a list. The heap is ordered according to the criterion $i \geq j \iff f_i t_i \geq f_j t_j$. Nominally, using a heap increases the cost of the main loop from $O(hm)$ to $O(hm \log m)$ because the heap is rebalanced when each t_j -value is updated. However, typical problem instances do not achieve this upper bound because the order of the $f_j t_j$ -values changes only slightly between iterations. The cost of updating each t_j -value can be reduced to unit time using a Fibonacci heap (Fredman and Tarjan 1987), yielding the same overall computation time.

3.4 Properties of the optimal portfolios

The nestedness property implies that Alma's expected utility is a discretely concave function of h .

Theorem 5 (Optimal portfolio valuation concave in h). *For $h = 2 \dots (m - 1)$,*

$$v(\mathcal{X}_h) - v(\mathcal{X}_{h-1}) \geq v(\mathcal{X}_{h+1}) - v(\mathcal{X}_h). \quad (19)$$

Proof. We will prove the equivalent expression $2v(\mathcal{X}_h) \geq v(\mathcal{X}_{h+1}) + v(\mathcal{X}_{h-1})$. Applying Theorem 3, we write $\mathcal{X}_h = \mathcal{X}_{h-1} \cup \{j\}$ and $\mathcal{X}_{h+1} = \mathcal{X}_{h-1} \cup \{j, k\}$. Define the random variables X_i as above. If $t_k \leq t_j$, then

$$\begin{aligned} 2v(\mathcal{X}_h) &= v(\mathcal{X}_{h-1} \cup \{j\}) + v(\mathcal{X}_{h-1} \cup \{j\}) \\ &\geq v(\mathcal{X}_{h-1} \cup \{k\}) + v(\mathcal{X}_{h-1} \cup \{j\}) \\ &= v(\mathcal{X}_{h-1} \cup \{k\}) + (1 - f_j)v(\mathcal{X}_{h-1}) + f_j \mathbb{E}[\max\{t_j, X_{h-1}\}] \\ &= v(\mathcal{X}_{h-1} \cup \{k\}) - f_j v(\mathcal{X}_{h-1}) + f_j \mathbb{E}[\max\{t_j, X_{h-1}\}] + v(\mathcal{X}_{h-1}) \\ &\geq v(\mathcal{X}_{h-1} \cup \{k\}) - f_j v(\mathcal{X}_{h-1} \cup \{k\}) + f_j \mathbb{E}[\max\{t_j, X_{h-1}\}] + v(\mathcal{X}_{h-1}) \\ &= (1 - f_j)v(\mathcal{X}_{h-1} \cup \{k\}) + f_j \mathbb{E}[\max\{t_j, X_{h-1}\}] + v(\mathcal{X}_{h-1}) \\ &= v(\mathcal{X}_{h-1} \cup \{j, k\}) + v(\mathcal{X}_{h-1}) \\ &= v(\mathcal{X}_{h+1}) + v(\mathcal{X}_{h-1}). \end{aligned} \quad (20)$$

The first inequality follows from the optimality of \mathcal{X}_h , while the second follows from the fact that adding k to \mathcal{X}_{h-1} can only increase its valuation.

If $t_k \geq t_j$, then the steps are analogous:

$$\begin{aligned} 2v(\mathcal{X}_h) &= v(\mathcal{X}_{h-1} \cup \{j\}) + v(\mathcal{X}_{h-1} \cup \{j\}) \\ &\geq v(\mathcal{X}_{h-1} \cup \{k\}) + v(\mathcal{X}_{h-1} \cup \{j\}) \\ &= (1 - f_k)v(\mathcal{X}_{h-1}) + f_k \mathbb{E}[\max\{t_k, X_{h-1}\}] + v(\mathcal{X}_{h-1} \cup \{j\}) \\ &= v(\mathcal{X}_{h-1}) - f_k v(\mathcal{X}_{h-1}) + f_k \mathbb{E}[\max\{t_k, X_{h-1}\}] + v(\mathcal{X}_{h-1} \cup \{j\}) \\ &\geq v(\mathcal{X}_{h-1}) - f_k v(\mathcal{X}_{h-1} \cup \{j\}) + f_k \mathbb{E}[\max\{t_k, X_{h-1}\}] + v(\mathcal{X}_{h-1} \cup \{j\}) \\ &= v(\mathcal{X}_{h-1}) + (1 - f_k)v(\mathcal{X}_{h-1} \cup \{j\}) + f_k \mathbb{E}[\max\{t_k, X_{h-1}\}] \\ &= v(\mathcal{X}_{h-1}) + v(\mathcal{X}_{h-1} \cup \{j, k\}) \\ &= v(\mathcal{X}_{h-1}) + v(\mathcal{X}_{h+1}) \end{aligned} \quad (21)$$

□

It follows that when \mathcal{X}_h is the optimal h -portfolio, for a given market, $v(\mathcal{X}_h)$ is $O(h)$. Example 2, in which $v(\mathcal{X}_h)$ can be made arbitrarily close to h , establishes the tightness of this bound.

4 Heterogeneous application costs

In this section, we turn to the more general problem in which the constant g_j represents the *cost* of applying to c_j and the student, whom we now call Ellis, has a *budget* of H to spend on

college applications. Applying Theorem 1, we assume $t_0 = 0$ and disregard c_0 throughout.

Definition 4 (Ellis's problem). Ellis's optimal college application portfolio is given by the solution to the following combinatorial optimization problem.

$$\begin{aligned} \text{maximize} \quad & v(\mathcal{X}) = \sum_{j \in \mathcal{X}} \left(t_j f_j \prod_{\substack{i \in \mathcal{X}: \\ i > j}} (1 - f_i) \right) \\ \text{subject to} \quad & \mathcal{X} \subseteq \mathcal{C}, \quad \sum_{j \in \mathcal{X}} g_j \leq H \end{aligned} \tag{22}$$

The optima for Ellis's problem are not necessarily nested, nor is the number of schools in the optimal portfolio necessarily increasing in H . For example, if $f = (0.5, 0.5, 0.5)$, $t = (1, 1, 219)$, and $g = (1, 1, 3)$, then it is evident that the optimal portfolio for $H = 2$ is $\{1, 2\}$ while that for $H = 3$ is $\{3\}$.

4.1 NP-completeness

In fact, Ellis's problem is NP-complete, as we will show by a transformation from the binary knapsack problem, which is known to be NP-complete (Garey and Johnson 1979).

Definition 5 (Decision form of knapsack problem). An *instance* consists of a set \mathcal{B} of m objects; utility values $u_j \in \mathbb{N}$ and weight $w_j \in \mathbb{N}$ for each $j \in \mathcal{B}$; and target utility $U \in \mathbb{N}$ and knapsack capacity $W \in \mathbb{N}$. The instance is called a *yes-instance* if and only if there exists a set $\mathcal{B}' \subseteq \mathcal{B}$ having $\sum_{j \in \mathcal{B}'} u_j \geq U$ and $\sum_{j \in \mathcal{B}'} w_j \leq W$.

Definition 6 (Decision form of Ellis's problem). An *instance* consists of an instance of Ellis's problem and a target valuation V . The instance is called a *yes-instance* if and only if there exists a portfolio $\mathcal{X} \subseteq \mathcal{C}$ having $v(\mathcal{X}) \geq V$ and $\sum_{j \in \mathcal{X}} g_j \leq H$.

Theorem 6. *The decision form of Ellis's problem is NP-complete.*

Proof. It is obvious that the problem is in NP.

Consider an instance of the knapsack problem, and we will construct an instance of Ellis's problem that is a yes-instance if and only if the corresponding knapsack instance is a yes-instance. Without loss of generality, we may assume that the objects in \mathcal{B} are indexed in increasing order of u_j , that each $u_j > 0$, and that the knapsack instance admits a feasible solution other than the empty set.

Let $U_{\max} = \sum_{j \in \mathcal{B}} u_j$ and $\delta = 1/mU_{\max} > 0$, and construct an instance of Ellis's problem with $\mathcal{C} = \mathcal{B}$, $H = W$, all $f_j = \delta$, and $t_j = u_j/\delta$ for all j . Clearly, $\mathcal{X} \subseteq \mathcal{C}$ is feasible for Ellis's problem if and only if it is feasible for the knapsack instance. Now, we observe that for any

nonempty \mathcal{X} ,

$$\begin{aligned}
\sum_{j \in \mathcal{X}} u_j &= \sum_{j \in \mathcal{X}} t_j f_j > \sum_{j \in \mathcal{X}} \left(t_j f_j \prod_{\substack{j' \in \mathcal{X}: \\ j' > j}} (1 - f_{j'}) \right) = v(\mathcal{X}) \\
&= \sum_{j \in \mathcal{X}} \left(u_j \prod_{\substack{j' \in \mathcal{X}: \\ j' > j}} (1 - \delta) \right) \geq (1 - \delta)^m \sum_{j \in \mathcal{X}} u_j \\
&\geq (1 - m\delta) \sum_{j \in \mathcal{X}} u_j \geq \sum_{j \in \mathcal{X}} u_j - m\delta U_{\max} = \sum_{j \in \mathcal{X}} u_j - 1.
\end{aligned} \tag{23}$$

This means that the value of an application portfolio \mathcal{X} for the corresponding knapsack instance is the smallest integer greater than $v(\mathcal{X})$. That is, $\sum_{j \in \mathcal{X}} u_j \geq U$ if and only if $v(\mathcal{X}) \geq U - 1$. Taking $V = U - 1$ completes the transformation and concludes the proof. \square

An intuitive extension of the greedy algorithm for Alma's problem is to iteratively add to \mathcal{X} the school k for which $[v(\mathcal{X} \cup \{k\}) - v(\mathcal{X})]/g_k$ is largest. However, the construction above shows that the objective function of Ellis's problem can approximate that of a knapsack problem with arbitrary precision. Therefore, in pathological examples such as the following, the greedy algorithm can achieve an arbitrarily poor approximation ratio.

Example 3. Let $t = (10, 2021)$, $f = (1, 1)$, $g = (1, 500)$, and $H = 500$. Then the greedy approximation algorithm produces the clearly inoptimal solution $\mathcal{X} = \{1\}$.

4.2 Pseudopolynomial-time dynamic program

In this subsection, we assume, with a small loss of generality, that $g_j \in \mathbb{N}$ for $j = 1 \dots m$ and $H \in \mathbb{N}$, and provide an algorithmic solution to Ellis's problem that runs in $O(Hm + m \log m)$ time and $O(Hm)$ space. The algorithm resembles a familiar dynamic programming algorithm for the binary knapsack problem (Dantzig 1957; *Wikipedia*, s.v. "Knapsack problem"). Because we cannot assume that $H \leq m$ (as was the case in Alma's problem), this represents a pseudopolynomial-time solution (Garey and Johnson 1979).

For $j = 0 \dots m$ and $h = 0 \dots H$, let $\mathcal{X}[j, h]$ denote the optimal portfolio using only the schools $\{1, \dots, j\}$ and costing no more than h , and let $V[j, h] = v(\mathcal{X}[j, h])$. It is clear that if $j = 0$ or $h = 0$, then $\mathcal{X}[j, h] = \emptyset$ and $V[j, h] = 0$. For convenience, we also define $V[j, h] = -\infty$ for all $h < 0$.

For the remaining indices, $\mathcal{X}[j, h]$ either contains j or not. If it does not contain j , then $\mathcal{X}[j, h] = \mathcal{X}[j-1, h]$. On the other hand, if $\mathcal{X}[j, h]$ contains j , then its value is $(1 - f_j)v(\mathcal{X}[j, h] \setminus \{j\}) + f_j t_j$. This requires that $\mathcal{X}[j, h] \setminus \{j\}$ make optimal use of the remaining budget over the remaining schools; that is, $\mathcal{X}[j, h] = \mathcal{X}[j-1, h - g_j] \cup \{j\}$. From these observations, we obtain the following Bellman equation for $j = 1 \dots m$ and $h = 1 \dots H$:

$$V[j, h] = \max\{V[j-1, h], (1 - f_j)V[j-1, h - g_j] + f_j t_j\} \tag{24}$$

with the convention that $-\infty \cdot 0 = -\infty$. The corresponding optimal portfolios can be computed by observing that $\mathcal{X}[j, h]$ contains j if and only if $V[j, h] > V[j-1, h]$. The optimal solution is

given by $\mathcal{X}[m, H]$. The algorithm below performs these computations and outputs the optimal portfolio \mathcal{X} .

Algorithm 2: Dynamic program for Ellis's problem with integral application costs.

Data: Utility values $t \in [0, \infty)^m$, admissions probabilities $f \in [0, 1]^m$, application costs $g \in \mathbb{N}^m$, budget $H \in \mathbb{N}$.

```

1 Index schools in ascending order by  $t$ ;
2 Fill a lookup table with the entries of  $V[j, h]$ ;
3  $h \leftarrow H$ ;
4  $\mathcal{X} \leftarrow \emptyset$ ;
5 for  $j = m, m - 1, \dots, 1$  do
6   if  $V[j - 1, h] < V[j, h]$  then
7      $\mathcal{X} \leftarrow \mathcal{X} \cup \{j\}$ ;
8      $h \leftarrow h - g_j$ ;
9   end
10 end
11 return  $\mathcal{X}$ 

```

Theorem 7 (Validity of Algorithm 2). *Algorithm 2 produces an optimal application portfolio for Ellis's problem in $O(Hm + m \log m)$ time and $O(Hm)$ space.*

Proof. Optimality follows from the foregoing discussion. Sorting t is $O(m \log m)$. The bottleneck step is the creation of the lookup table for $V[j, h]$ in line 2. Each entry is generated in unit time, and the size of the table is $O(Hm)$. \square

4.3 Fully polynomial-time approximation scheme

As with the knapsack problem, Ellis's problem admits a complementary dynamic program that iterates on the value of the cheapest portfolio instead of on the cost of the most valuable portfolio. We will use this algorithm as the basis for a fully polynomial-time approximation scheme for Ellis's problem that uses $O(m^3/\varepsilon)$ time and space. Here we assume, with a small loss of generality, that each t_j is a natural number.

We will represent the approximate value of portfolios using a fixed-point decimal with a precision of P , where P is the number of digits to retain after the decimal point. Let $r[x] = \lfloor 10^P x \rfloor 10^{-P}$ denote the value of x rounded down to its nearest fixed-point representation. Since $\bar{U} = \sum_{j \in \mathcal{C}} f_j t_j$ is an upper bound on the value of any portfolio, and since we will ensure that each fixed-point approximation is an underestimate of the portfolio's true value, the set \mathcal{V} of possible valuations possible in the fixed-point framework is finite:

$$\mathcal{V} = \left\{ 0, 1 \times 10^{-P}, 2 \times 10^{-P}, \dots, r[\bar{U} - 1 \times 10^{-P}], r[\bar{U}] \right\} \quad (25)$$

Then $|\mathcal{V}| = \bar{U} \times 10^P + 1$.

For the remainder of this subsection, unless otherwise specified, the word *value* refers to a portfolio's value within the fixed-point framework, with the understanding that this is an approximation. We will account for the approximation error below when we prove the dynamic program's validity.

For integers $0 \leq j \leq m$ and $v \in [-\infty, 0) \cup \mathcal{V}$, let $\mathcal{W}[j, v]$ denote the least expensive portfolio that uses only schools $\{1, \dots, j\}$ and has value at least v , if such a portfolio exists. Denote its value by $G[j, v]$, where $G[j, v] = \infty$ if $\mathcal{W}[j, v]$ does not exist. It is clear that if $v \leq 0$, then $\mathcal{W}[j, v] = \emptyset$ and $G[j, v] = 0$, and that if $j = 0$ and $v > 0$, then $G[j, v] = \infty$. For the remaining indices (where $j, v > 0$), we claim that

$$G[j, v] = \begin{cases} \infty, & t_j < v \\ \min\{G[j-1, v], g_j + G[j-1, v - \Delta_j(v)]\}, & t_j \geq v \end{cases} \quad (26)$$

$$\text{where } \Delta_j(v) = \begin{cases} r \left[\frac{f_j}{1-f_j} (t_j - v) \right], & f_j < 1 \\ \infty, & f_j = 1 \end{cases} \quad (27)$$

In the $t_j < v$ case, any feasible portfolio must be composed of schools with utility less than v , and therefore its valuation can not equal v , meaning that $\mathcal{W}[j, v]$ is undefined. In the $t_j \geq v$ case, the first argument to $\min\{\}$ says simply that omitting j and choosing $\mathcal{W}[j-1, v]$ is a permissible choice for $\mathcal{W}[j, v]$. If, on the other hand, $j \in \mathcal{W}[j, v]$, then

$$v(\mathcal{W}[j, v]) = (1 - f_j)v(\mathcal{W}[j, v] \setminus \{j\}) + f_j t_j. \quad (28)$$

Therefore, the subportfolio $\mathcal{W}[j, v] \setminus \{j\}$ must have a value of at least $v - \Delta$, where Δ satisfies $v = (1 - f_j)(v - \Delta) + f_j t_j$. When $f_j < 1$, the solution to this equation is $\Delta = \frac{f_j}{1-f_j} (t_j - v)$. By rounding this value down, we ensure that the true valuation of $\mathcal{W}[j, v]$ is *at least* $v - \Delta$.

When $f_j = 1$, then the singleton $\{j\}$ has $v(\{j\}) \geq v$, so

$$G[j, v] = \min\{G[j-1, v], g_j\}. \quad (29)$$

Defining $\Delta = \infty$ in this case ensures that $G[j-1, v - \infty] = 0$ as required.

Once $G[j, v]$ has been calculated at each index, the associated portfolio can be found by applying the observation that $\mathcal{W}[j, v]$ contains j if and only if $G[j, v] < G[j-1, v]$. Then an approximate solution to Ellis's problem is obtained by computing the optimal objective value

$\max\{w : G[m, w] \leq H\}$ and corresponding portfolio.

Algorithm 3: Fully polynomial-time approximation scheme for Ellis's problem.

Data: Utility values $t \in \mathbb{N}^m$, admissions probabilities $f \in (0, 1]^m$, application costs $g \in (0, \infty)^m$, budget $H \in (0, \infty)^m$, tolerance $\varepsilon \in (0, 1)$.

```

1 Index schools in ascending order by  $t$ ;
2 Set precision  $P \leftarrow \lceil \log_{10}(m^2/\varepsilon\bar{U}) \rceil$ ;
3 Fill a lookup table with the entries of  $G[j, h]$ ;
4  $v \leftarrow \max\{w \in \mathcal{V} : G[m, w] \leq H\}$ ;
5  $\mathcal{X} \leftarrow \emptyset$ ;
6 for  $j = m, m-1, \dots, 1$  do
7   if  $G[j, v] < \infty$  and  $G[j, v] < G[j-1, v]$  then
8      $\mathcal{X} \leftarrow \mathcal{X} \cup \{j\}$ ;
9      $v \leftarrow v - \Delta_j(v)$ ;
10  end
11 end
12 return  $\mathcal{X}$ 

```

Theorem 8 (Validity of Algorithm 3). *Algorithm 3 produces a $(1 - \varepsilon)$ -optimal application portfolio for Ellis's problem in $O(m^3/\varepsilon)$ time.*

Proof. (Optimality.) Let \mathcal{W} denote the output of Algorithm 3 and \mathcal{X} the true optimum. We know that $v(\mathcal{X}) \leq \bar{U}$, and because each singleton portfolio is feasible, \mathcal{X} must be more valuable than the average singleton portfolio; that is, $v(\mathcal{X}) \geq \bar{U}/m$.

Because $\Delta_j(v)$ is rounded down in the recursion relation defined by (26) and (27), if $j \in \mathcal{W}[j, v]$, then the true value of $(1 - f_j)v(\mathcal{W}[j-1, v - \Delta_j(v)]) + f_j t_j$ may exceed the fixed-point value v of $\mathcal{W}[j, v]$, but not by more than 10^{-P} . This error accumulates with each school added to \mathcal{W} , but the number of additions is at most m . Therefore, where $v'(\mathcal{W})$ denotes the fixed-point valuation of \mathcal{W} recorded in line 4 of the algorithm, $v(\mathcal{W}) - v'(\mathcal{W}) \leq m10^{-P}$.

We can define $v'(\mathcal{X})$ analogously as the fixed-point valuation of \mathcal{X} when its elements are added in index order and its value is updated and rounded down to the nearest multiple of 10^{-P} at each addition in accordance with (28). By the same logic, $v(\mathcal{X}) - v'(\mathcal{X}) \leq m10^{-P}$. The optimality of \mathcal{W} in the fixed-point environment implies that $v'(\mathcal{W}) \geq v'(\mathcal{X})$.

Applying these observations, we have

$$v(\mathcal{W}) \geq v'(\mathcal{W}) \geq v'(\mathcal{X}) \geq v(\mathcal{X}) - m10^{-P} \geq \left(1 - \frac{m^2 10^{-P}}{\bar{U}}\right) v(\mathcal{X}) \geq (1 - \varepsilon) v(\mathcal{X}) \quad (30)$$

which establishes the approximation bound.

(Computational resources.) The bottleneck step is the creation of the lookup table in line 3, whose size is $O(|\mathcal{V}|m)$. Since

$$|\mathcal{V}| = \bar{U} \times 10^P + 1 = \bar{U} \times 10^{\lceil \log_{10}(m^2/\varepsilon\bar{U}) \rceil} + 1 \leq \frac{m^2}{\varepsilon} \times \text{const.} \quad (31)$$

is $O(m^2/\varepsilon)$, the time complexity is as promised. \square

Since these bounds are polynomial in m and $1/\varepsilon$, Algorithm 3 is a fully polynomial-time approximation scheme for Ellis’s problem (Vazirani 2001).

Algorithms 2 and 3 can be written using recursive functions instead; however, since each function references itself *twice*, the function values at each index must be recorded in a lookup table or otherwise memoized to prevent an exponential number of calls from forming on the stack.

5 (WIP) Numerical experiments

In this section, we present the results of numerical experiments designed to confirm the time complexity results established above. In both experiments, markets were generated by drawing t_j independently from an exponential distribution with a scale parameter of ten and rounding up to the nearest integer. To achieve partial negative correlation between t_j and f_j , we then set $f_j = 1/(t_j + 10 + Q)$, where Q is a uniform draw from interval $[0, 1)$. In Experiment 1, which concerns Algorithm 1, we take each $g_j = 1$ and set $H = h = \lfloor m/2 \rfloor$. In Experiment 2, which concerns Algorithms 2 and 3, each g_j is drawn uniformly from the set $\{5, \dots, 10\}$ and we take $H = \lfloor \frac{1}{2} \sum g_j \rfloor$. At each combination of the experimental variables, we generated 50 markets, and each computation was repeated three times, with the fastest of the three recorded as the computation time. Therefore, each cell of each table represents 150 computations. We report the mean and standard deviation across the 50 markets. Where applicable, we do not count the time required to sort the entries of t .

The dynamic programs were implemented using recursive functions and memoization. Our implementation of Algorithm 3 differed from the text in that portfolio valuations were recorded as fixed-point *binary* numbers instead, with the definitions of P and \mathcal{V} modified accordingly.

6 References

- Budish, Eric. 2011. “The Combinatorial Assignment Problem: Approximate Competitive Equilibrium from Equal Incomes.” *Journal of Political Economy* 119 (6): 1061–1103. <https://doi.org/10.1086/664613>.
- Dantzig, George B. 1957. “Discrete-Variable Extremum Problems.” *Operations Research* 5 (2): 266–88.
- Fisher, Marshall, George Nemhauser, and Laurence Wolsey. 1978. “An analysis of approximations for maximizing submodular set functions—I.” *Mathematical Programming* 14: 265–94.
- Fredman, Michael Lawrence and Robert Tarjan. 1987. “Fibonacci heaps and their uses in improved network optimization algorithms.” *Journal of the Association for Computing Machinery* 34 (3): 596–615.
- Fu, Chao. 2014. “Equilibrium Tuition, Applications, Admissions, and Enrollment in the College Market.” *Journal of Political Economy* 122 (2): 225–81. <https://doi.org/10.1086/675503>.
- Garey, Michael and David Johnson. 1979. *Computers and Intractability: A Guide to the Theory of NP-Completeness*. New York: W. H. Freeman and Company.
- Meucci, Attilio. 2005. *Risk and Asset Allocation*. Berlin: Springer-Verlag, 2005.
- Othman, Abraham, Eric Budish, and Tuomas Sandholm. 2010. “Finding Approximate Competitive Equilibria: Efficient and Fair Course Allocation.” In *Proceedings of 9th International Conference on Autonomous Agents and Multiagent Systems*. New York: ACM. <https://dl.acm.org/doi/abs/10.5555/1838206.1838323>.
- Roazanov, Mark and Arie Tamir. 2020. “The nestedness property of the convex ordered median location problem on a tree.” *Discrete Optimization* 36: 100581. <https://doi.org/10.1016/j.disopt.2020.100581>.
- Vazirani, Vijay. 2001. *Approximation Algorithms*. Berlin: Springer.